

A Message to Watermarkers

Abstract

Many AI labs including but not limited to Anthropic, OpenAI, Google DeepMind, apply watermarks globally onto outputs produced by their models. Project OffQuill aims to send a message, one backed by the open source community's actions against watermarking such as the **watermarks-remover** GitHub repository. We will also cover the aims and motives behind project OffQuill, while briefly touching on how it was developed.

1. Introduction

AI labs watermark output. So what? Everyday people don't really care. Sure, a student might worry that their teacher can now see that they are cheating on homework, or a teacher might be rejoicing at their ability to see if students are using AI to do assignments. But the open source community has reacted differently, and for good reasons. To understand these reasons, we need to take a closer look at the meaning of open source.

Now, easily enough you can Google it and find that the Merriam-Webster Dictionary defines open source as "software having freely available source code for possible modification and redistribution, as well as being publicly available for general community use." In simple terms this means code that is shared publicly so anyone can look at it, change it, and use it. The rule of thumb core open source values are listed below:

- Transparency - The source code itself must be completely visible to everyone.
- Collaboration - No single individual or company can solve a problem all by themselves.
- Innovation - The philosophy embraces a "release early and release often" mindset.
- Meritocracy - The best ideas win through peer review and active contribution.

If you pay attention and piece together the puzzle, you'll start to realise why the open source community opposes AI output watermarking so fiercely. A watermark is an identifier that a user didn't ask for being added into their model output.

Detection requires the lab's key and infrastructure, again relating to transparency and arguably meritocracy as the lab's authority is not completely peer reviewable. The ability to collaborate and innovate freely is also restricted. For greater depth, allow me to provide some example scenarios below:

1. Firstly, transparency. AI labs restrict public access to the key behind these watermarking algorithms, keeping them tightly guarded and only accessible to highly privileged organisations. This plainly violates the open source value.
2. Second is collaboration. Watermarking forces open source developers to either risk using text with unknown alterations or exclude model generated content entirely. Because detecting these watermarks requires the lab's private infrastructure, this restriction allows a single party to dictate terms, directly violating the cooperative culture of the open-source ecosystem.
3. Next is innovation. AI labs use special algorithms and a key to the watermarking pattern. (We'll cover more on this later) However, by forcing a model to alter its statistical choices just to fit with this pattern, this can cause degradation in the models output quality and reasoning. This violates the open source commitment to maximum innovation and technological advancement.
4. Last is meritocracy. Imagine you're on a contract with a company that needs software, and you're almost done. You're just a regular developer, so you hand the code to AI for a thorough security audit and bug fix. However, the bug fixes contain watermarks. The company checks the code, finds traces and refuses to pay you in full, thinking you don't actually know how to code, you're just using AI. But this wasn't true. You did 90% of the work and only handed it to AI for the finishing touches. But now people don't trust you, all because an AI watermarked your work when you asked it for light help.

Do you better understand the hidden toll of these watermarks on the open source community now? All of this for what, to ensure that we can trace artificial content that spreads misinformation? This is not only mathematically flawed, but is actively harmful to innovation and freedom of speech. What do I mean by mathematically flawed? Ask any mathematician with a brain and knowledge on the context, they'll tell you that it's impossible, because you cannot prove a statistical certainty from a flexible and limited sample.

2. Timeline of events

Let's take a closer look at the timeline of events related to the open source community and AI watermarking.

In late 2023, AI labs began experimenting with the first methods of watermarking, more specifically the SynthID system.

In May 2026 at the Google I/O conference, the massive scale of closed watermarking was revealed as over 100 billion images and videos have been watermarked using SynthID.

In July 2026, tech companies scrambled to begin preparing for inevitable European regulations.

August 11th, 2026. The date that shook the world. Anthropic announced that it would globally apply watermarks to all of its Claude model outputs.

August 14th 2026. Within hours of Anthropic confirming its global watermark implementation, open source developer Guillaume Meyer develops and publishes an open source bypass.

August 15th 2026. The GitHub repository by Meyer, titled **watermarks-remover**, goes viral globally in a matter of hours, being awarded thousands of GitHub stars and attracting open source contributors.

Late August 2026. The open source community has rapidly moved as a whole, and has retaliated to the watermarks by releasing countless tools designed to thoroughly scrub them from not just Anthropic, but any AI model output.

3. OffQuill's role

We have discussed watermarking, open source values, and the timeline, but I have not yet explained what OffQuill is. Put simply, OffQuill is an offline, local, lighter private version of the original watermarks-remover repository that was designed with a hacker's mindset. Now is the moment in which I will touch on the development process of Project Feather. (OffQuill)

What makes OffQuill different?

Modern watermark removers rely on 2 layers, static word and or phrase replacement, invisible character stripping as well as sending the text to an AI via API that doesn't watermark, such as GLM or Deepseek for thorough paraphrasing.

OffQuill borrows the first static layer emphasizing on invisible character stripping and light phrase swapping, but takes a unique approach with the LLM paraphrasing and adds a special optional touch.

Instead of relying on cloud LLM APIs which may retain your data, it uses the built in Gemini Nano local AI model built into Chrome by Google, which excels at paraphrasing to reword the text into sounding more human first, stripping it of any cold, robotic voice. (Finally a good use for the on device AI feature)

However, I realised that Gemini also applies its own watermarks, meaning even if the text passed ZeroGPT, QuillBot, or GPTZero, it was still vulnerable to Google's own scanners. To overcome this and not just break, but completely shatter the mathematical watermark, OffQuill locally translates your text from English to Japanese, and then Japanese back to English using RBMT. (Rule Based Machine Translation). Because wild variations in sentence structure between these two languages, translating them both ways transforms the token distribution using a non AI based translator, effectively scrambling the mathematical signature beyond recovery while still retaining meaning of the text. (Is this not innovation?)

In summary, OffQuill works for Chrome only, but allows users to privately humanise text and code (beta feature for code humanising) locally. It is designed to be cross OS, and easy to use. OffQuill is not meant to compete with watermarks-remover, but accompany it for users who prefer privacy and no setup.

4. Counter arguments against closed source defense

For all those who oppose the removal of watermarks, please find below several points that I hope will change your minds on the matter:

Bad actors trying to spread election disinformation or targeted scams will not politely use watermarked models from companies like Google or Anthropic.

Anyone with a decent computer can download a powerful, completely open source Large Language Model (like Meta's Llama) and strip the watermarking code out entirely on their own.

For text watermarking to work, it has to rely on specific word choices or statistical patterns. The open-source community points out that a bad actor can easily strip a watermark by running the text through a basic paraphrasing tool, changing a few words manually, or translating it back and forth between languages just like OffQuill.

If the public is taught to believe "No watermark means human written and trusted," they will be easily fooled by malicious AI text that simply had its watermark erased.

Watermarks cannot distinguish between a bot farm generating fake news and a non native English speaker using AI to proofread a resume or email.

Open-source argues that watermarking unfairly penalizes students, programmers, and professionals who use AI as a collaborative tool, as their authentic work gets falsely flagged with a "suspicious" AI stamp.

If a platform wants to check text for a watermark, they have to send that text to a corporate API. The open source community argues this compromises user privacy by requiring content to be scanned by Big Tech.

5. Final words

As this document nears its end, and I must go to bed soon, I would like to end with a few words.

Firstly, please do not think of me as anti public safety or pro AI slop. I genuinely agree with watermarking AI images and videos, as I believe creativity and art should only be done authentically by humans. However, the reason I feel differently about text is because as an open source developer, using AI is the only way we can keep with Silicon Valley as we do not possess the vast resources that they do. Additionally, text arguably has more legitimate uses than AI anime pictures or talking fruit videos.

Secondly, all information in this document is provided “as is”, sort of like the MIT license. I am not a professional researcher, but have tried my best to stick to the truth in this document. If there are any factual inaccuracies, I offer my sincerest apologies to all readers and assure you that it was not intentional. Additionally, I do not seek to force the opinions mentioned in this document on anyone. This paper was written to accompany the OffQuill GitHub release, not start a global debate.

Finally, I would like to end by mentioning the following entities for their contributions in the modern tech world:

- All of the Zhipu, Deepseek, MiMo, Moonshot and Qwen team for open source modern AI.
- Guillaume Meyer for sparking the open source retaliation against closed watermarking by Big Tech.